

CS 481

***Artificial Intelligence Language
Understanding***

March 28, 2023

Announcements / Reminders

- Please follow the Week 20 To Do List instructions
- Written Assignment #03 is due on ~~Sunday 03/26/23~~
TONIGHT at 11:59 PM CST
- Programming Assignment #02 is due on Sunday
04/02/23 at 11:59 PM CST
- Final Exam date:
Thursday 04/27/2023 (last week of classes!)
 - Ignore the date provided by the Registrar
 - Section 02 [Online]: contact Mr. Charles Scott
(scott@iit.edu) to arrange your exam

Plan for Today

- **Sentiment Analysis**
- **Words and their meaning**

Challenge

- **We know word relationships exist**
- **How can we quantify them in a automated fashion?**
- **How do we represent them in numerical way?**
- **How can we use them in computational models and processes?**

Computational Models of Meaning

"a word is characterized by the company it keeps"

- John Rupert Firth (English linguist)

"In most cases, the meaning of a word is its use"

- Ludwig Wittgenstein (Austrian philosopher)

"If A and B have almost identical environments we say that they are synonyms."

- Zellig Harris (American linguist)

Words + Their Environment: Example

- **Suppose you see these sentences:**

- *Ong choi is delicious sautéed with garlic.*
- *Ong choi is superb over rice*
- *Ong choi leaves with salty sauces*

- **And you've also seen these:**

- *...spinach sautéed with garlic over rice*
- *Chard stems and leaves are delicious*
- *Collard greens and other salty leafy greens*

- **Conclusion:**

- **Ong choi is a leafy green like spinach, chard, or collard greens**
- **We could conclude this based on words like "*leaves*" and "*delicious*" and "*sautéed*"**

Computational Models of Meaning

- So:
 - words are defined by their environments (the words around them)
- How can we represent word meaning with word environment?
 - **Vector semantics**

Vector Semantics: Two Ideas

- Idea 1:

- Let's define the **meaning of a word by its distribution in language use** (neighboring words or grammatical environments)

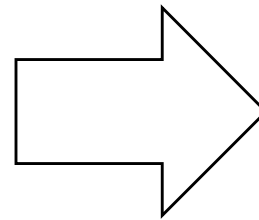
- Idea 2:

- Let's define the **meaning of a word as a point in space**

Bag of Words: Strings Representation

Some document:

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



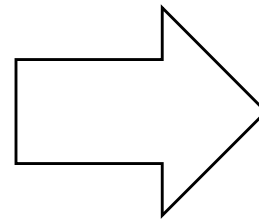
Word:	Frequency:
it	6
I	5
the	4
to	3
and	3
seen	2
yet	1
whimsical	1
times	1
....	...

Bag of words assumption: word/token position does not matter.

Bag of Words: Meaning Ignored!

Some document:

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!



Word:	Frequency:
it	6
I	5
the	4
to	3
and	3
seen	2
yet	1
whimsical	1
times	1
....	...

Bag of words assumption: word/token position does not matter.

Connotation as a Point in Space

- Words seem to vary along **three affective DIMENSIONS**:
 - valence**: the pleasantness of the stimulus
 - arousal**: the intensity of emotion provoked by the stimulus
 - dominance**: the degree of control exerted by the stimulus

	Word	Score		Word	Score
valence	love	1.000		toxic	0.008
	happy	1.000		nightmare	0.005
arousal	elated	0.960		mellow	0.069
	frenzy	0.965		napping	0.046
dominance	powerful	0.991		weak	0.045
	leadership	0.983		empty	0.081

Source: NRC VAD Lexicon (<https://saifmohammad.com/WebPages/nrc-vad.html>)

Vector Semantics

- The idea:
 - represent a word as **a point in a multidimensional semantic space** that is **derived from the distributions of word neighbors**

Point in Space Based on Distribution

- Each **word = a vector**
 - not just "*good*" or "*word₄₅*"
- **Similar words: “nearby in semantic space”**
- We build this space automatically by seeing which words are nearby in text



Vector Semantics: Words as Vectors



Source: Signorelli, Camilo & Arsiwalla, Xerxes. (2019). *Moral Dilemmas for Artificial Intelligence: a position paper on an application of Compositional Quantum Cognition*

Word Embedding: Definition

Word Embedding:

*a term used for the **representation of words** for text analysis, typically in the form of a real-valued vector that encodes the meaning of the word such that the words that are closer in the vector space are expected to be similar in meaning*

from Wikipedia

Word Embedding

- **Embedding:**
 - “embedded into a space”
 - mapping from one space or structure to another
- The **standard way to represent meaning** in NLP
- Fine-grained **model of meaning for similarity**

The Why: Sentiment Analysis

- Using **words only**:

- a feature is a word identity

- for example

- feature $x_5 = \begin{cases} \textcolor{red}{1} & \text{if the previous word was 'terrible'} \\ \textcolor{green}{0} & \text{otherwise} \end{cases}$

- requires exact same word to be in training and test

The Why: Sentiment Analysis

- Using **embeddings**:
 - a feature is a word vector
 - the previous word was vector $[35, 22, 17]$
 - now in the test set we might see a similar vector $[34, 21, 14]$
 - we can **generalize** to **similar but unseen** words

Term-Document Matrix

- Each document is represented by a vector of words

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

Term-Document Matrix

- Vectors are similar for the two comedies
 - “As you like it” and “Twelfth Night”

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

- But comedies are different than the other two
 - more *fools* and *wit* and fewer *battles*

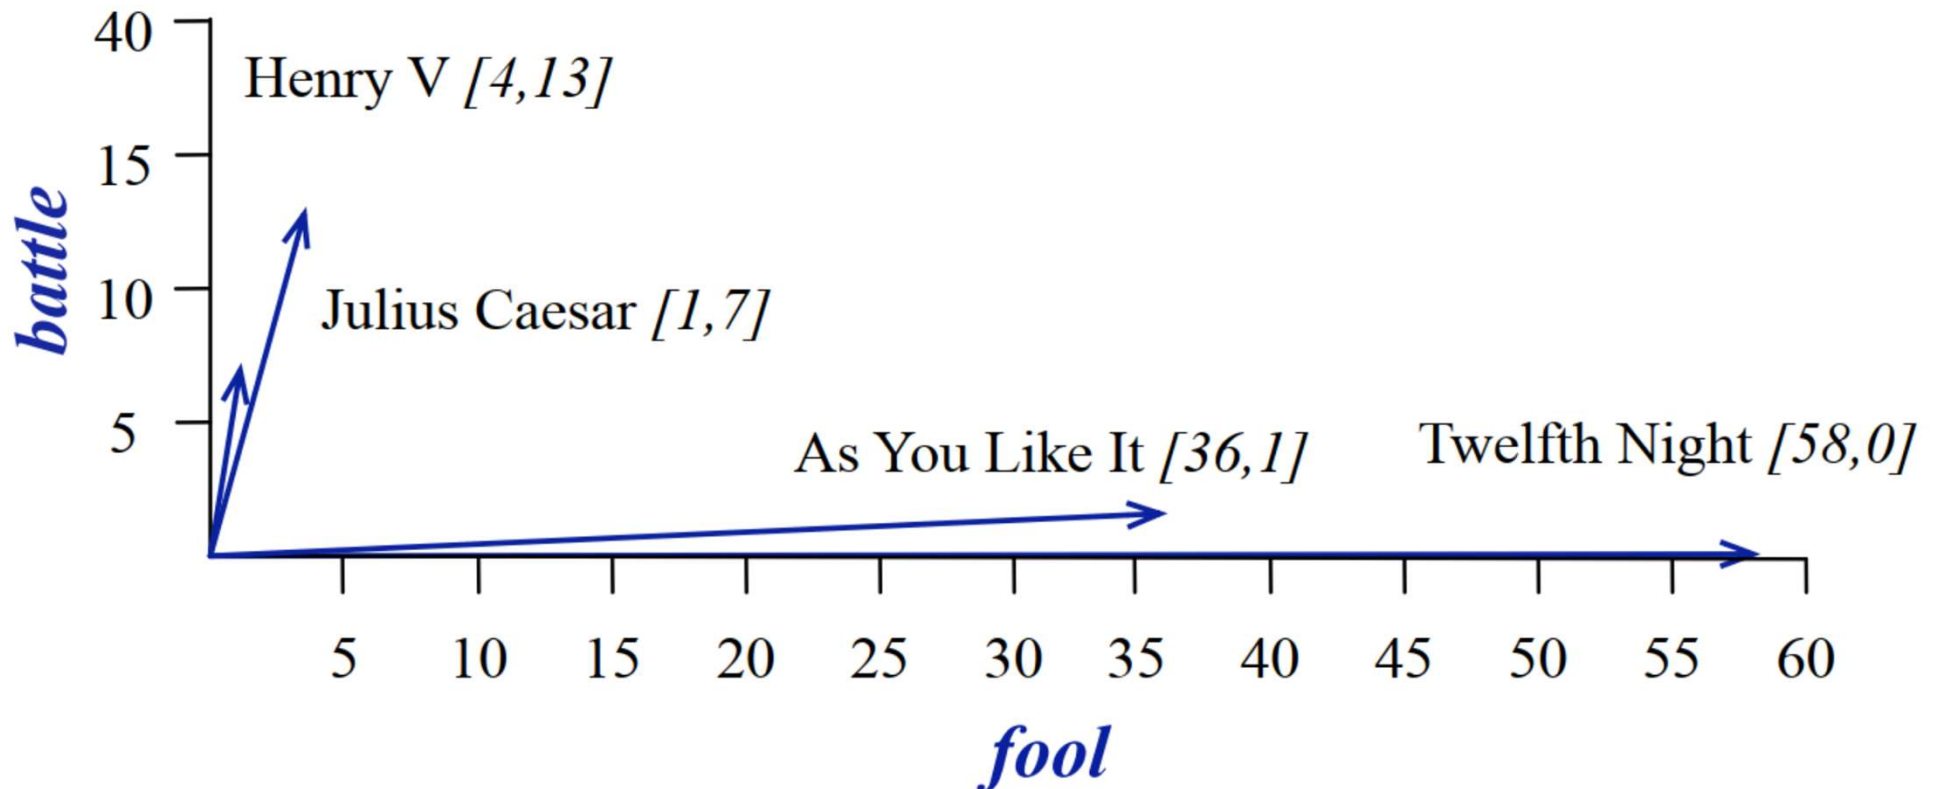
Term-Document Matrix

- Vectors are similar for the two comedies
 - “As you like it” and “Twelfth Night”

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

- But comedies are different than the other two
 - more *fools* and *wit* and fewer *battles*

Document Vector Visualization



Words as Vectors

- *battle* is "the kind of word that occurs in Julius Caesar and Henry V"

	As You Like It	Twelfth Night	Julius Caesar	Henry V
battle	1	0	7	13
good	114	80	62	89
fool	36	58	1	4
wit	20	15	2	3

- *fool* is "the kind of word that occurs in comedies, especially Twelfth Night"

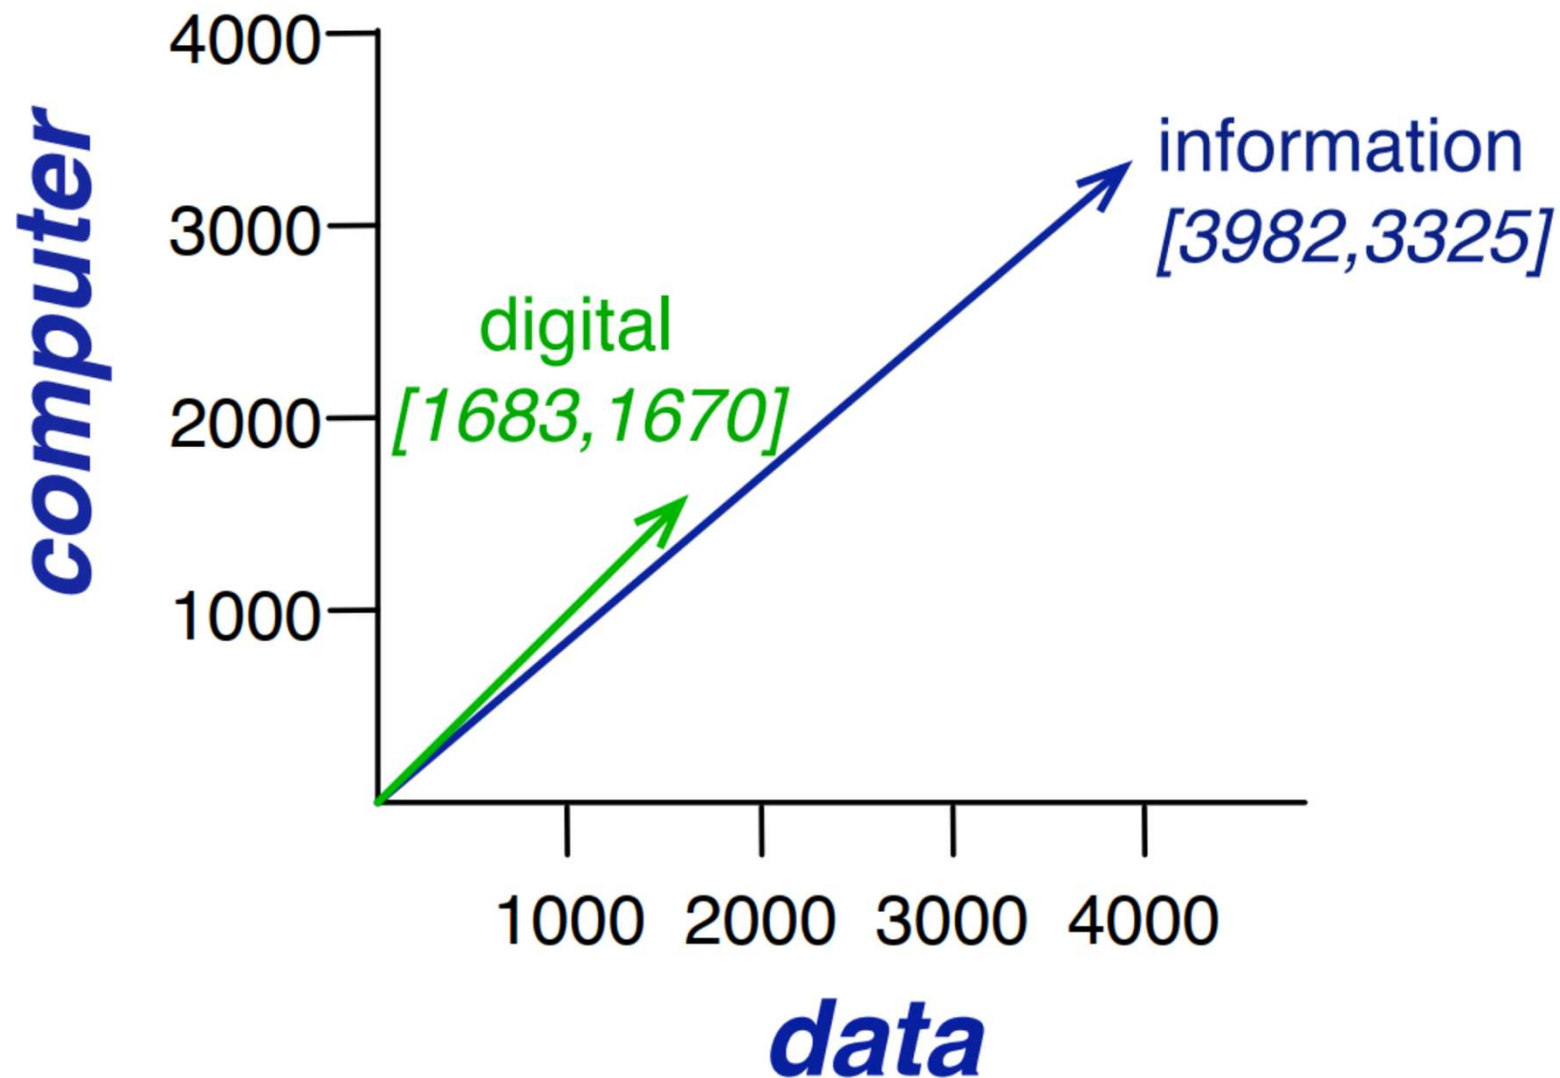
Word-Word (Term-Context) Matrix

- Two words are **similar in meaning** if their **context vectors are similar**

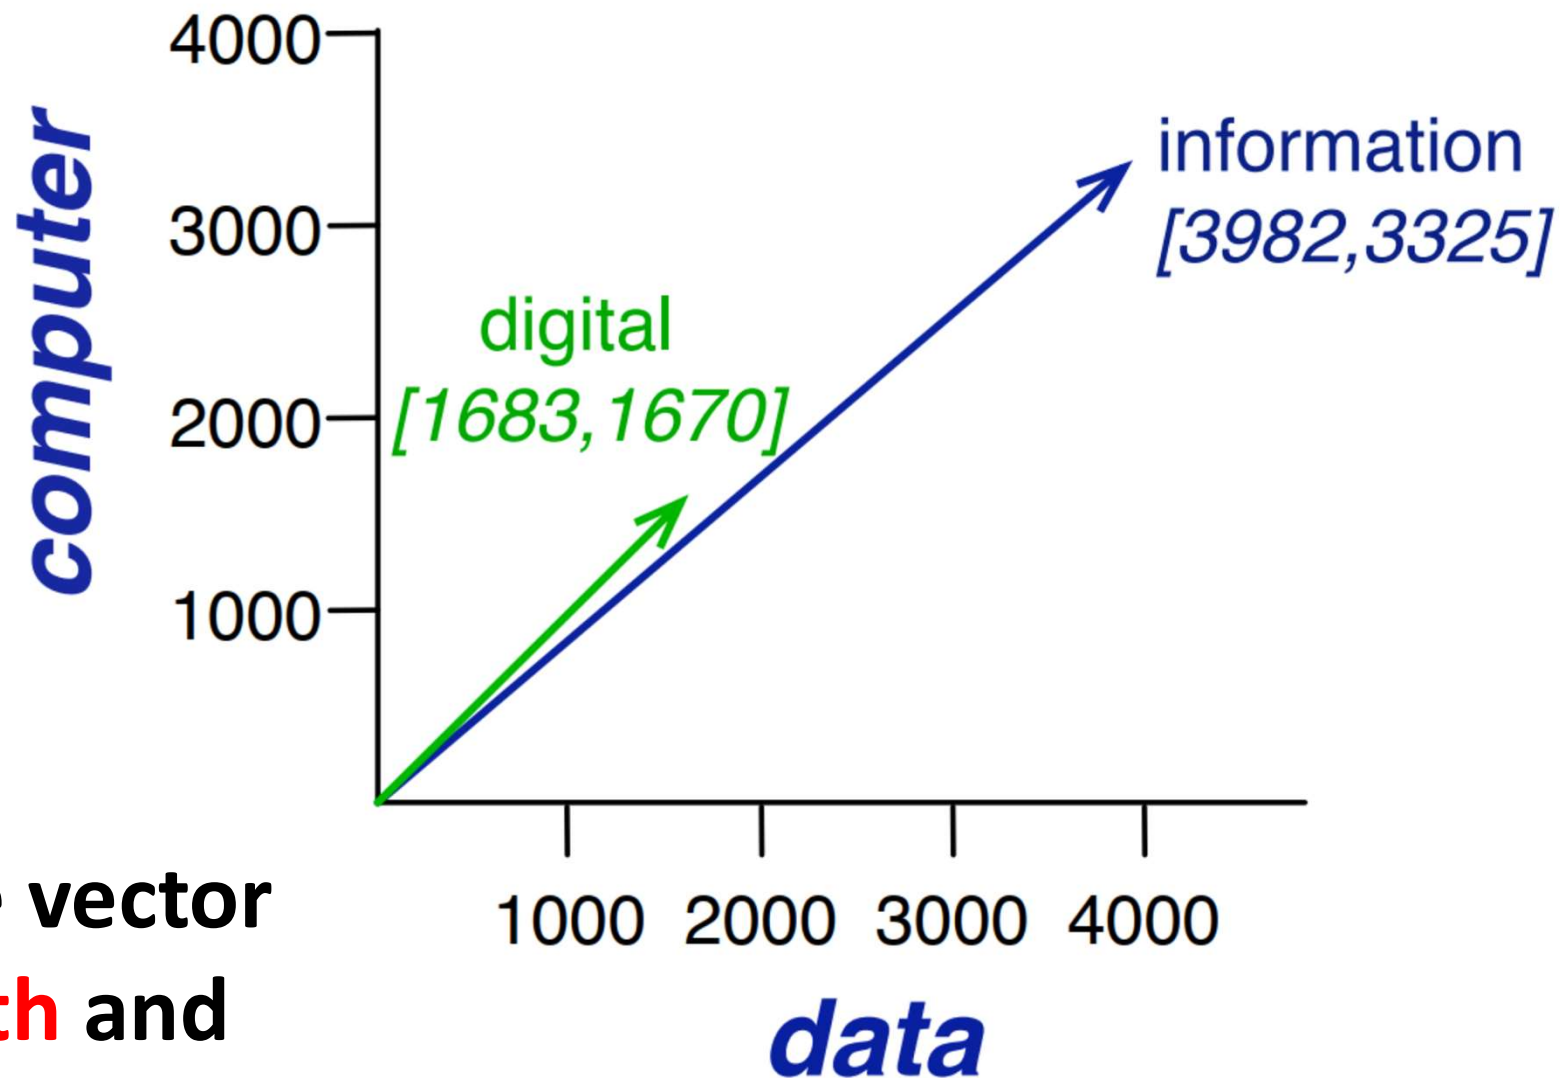
is traditionally followed by **cherry** pie, a traditional dessert
often mixed, such as **strawberry** rhubarb pie. Apple pie
computer peripherals and personal **digital** assistants. These devices usually
a computer. This includes **information** available on the internet

	aardvark	...	computer	data	result	pie	sugar	...
cherry	0	...	2	8	9	442	25	...
strawberry	0	...	0	0	1	60	19	...
digital	0	...	1670	1683	85	5	4	...
information	0	...	3325	3982	378	5	13	...

Document Vector Visualization



Document Vector Visualization



Note vector
length and
direction

Vector Dot / Scalar Product

Given two vectors \mathbf{a} and \mathbf{b} (N - vector space dimension):

$$\mathbf{a} = [a_1, a_2, \dots, a_N] \text{ and } \mathbf{b} = [b_1, b_2, \dots, b_N]$$

their vector dot/scalar product is:

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^N a_i * b_i = a_1 * b_1 + a_2 * b_2 + \dots + a_N * b_N$$

Using matrix representation:

$$\mathbf{a} \cdot \mathbf{b} = \mathbf{a} \mathbf{b}^T = [a_1 \quad a_2 \quad \dots \quad a_N] \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_N \end{bmatrix}$$

Vector Dot / Scalar Product

- Vector dot/scalar product is a **scalar**:

$$\mathbf{a} \cdot \mathbf{b} = \sum_{i=1}^N a_i * b_i = a_1 * b_1 + a_2 * b_2 + \dots + a_N * b_N$$

- Vector dot/scalar:
 - high values when the two vectors have large values in the same dimensions
 - **useful similarity measure**

Vector Dot / Scalar Product: Problem

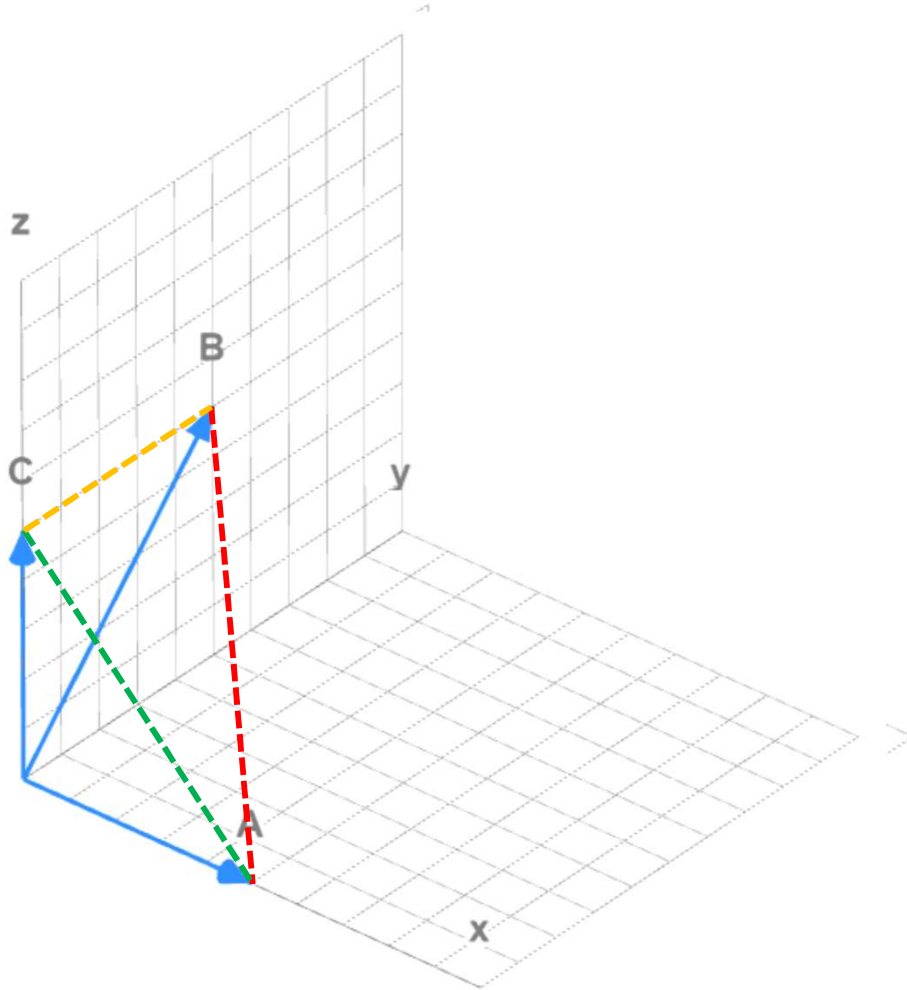
- Dot product **favors long vectors**: higher if a vector is longer (has higher values in many dimension)
- Vector length:

$$|\mathbf{a}| = \sqrt{\sum_{i=1}^N a_i^2}$$

- Frequent words (of, the, you) have long vectors (since they occur many times with other words).
- dot product overly **favors frequent words**

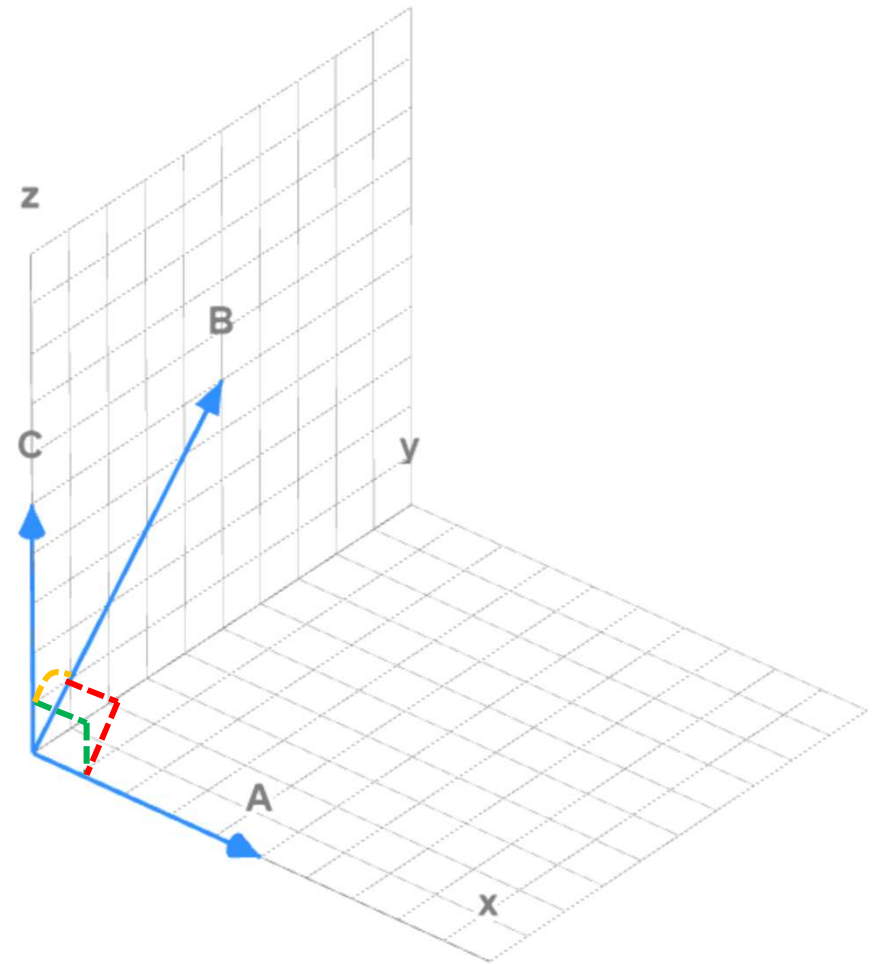
Alternative: Cosine Similarity

Euclidean distance



$$D(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

Cosine similarity



$$D(x, y) = \cos(\theta) = \frac{x \cdot y}{\|x\| \|y\|}$$

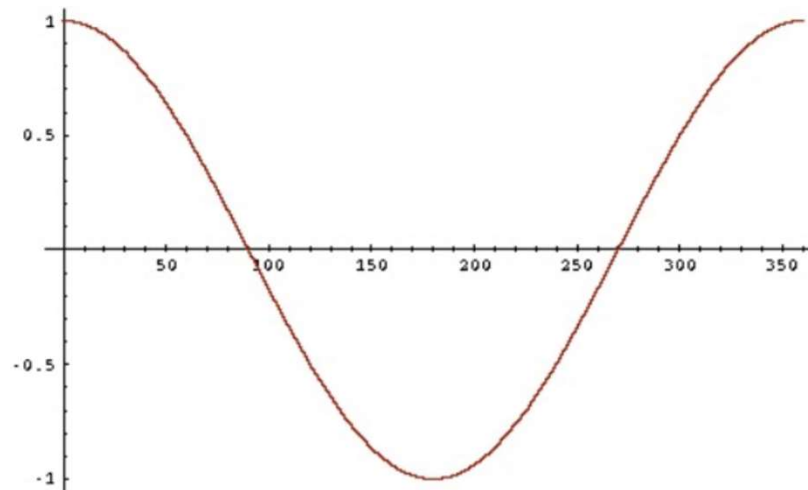
Word Similarity | Cosine Similarity

$$\text{cosine}(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

Where: v and w are two different word vectors

Word Similarity | Cosine Similarity

- **-1: vectors point in opposite directions**
- **+1: vectors point in same directions**
- **0: vectors are orthogonal**



But since raw frequency values are non-negative, the cosine for term-term matrix vectors ranges from 0–1

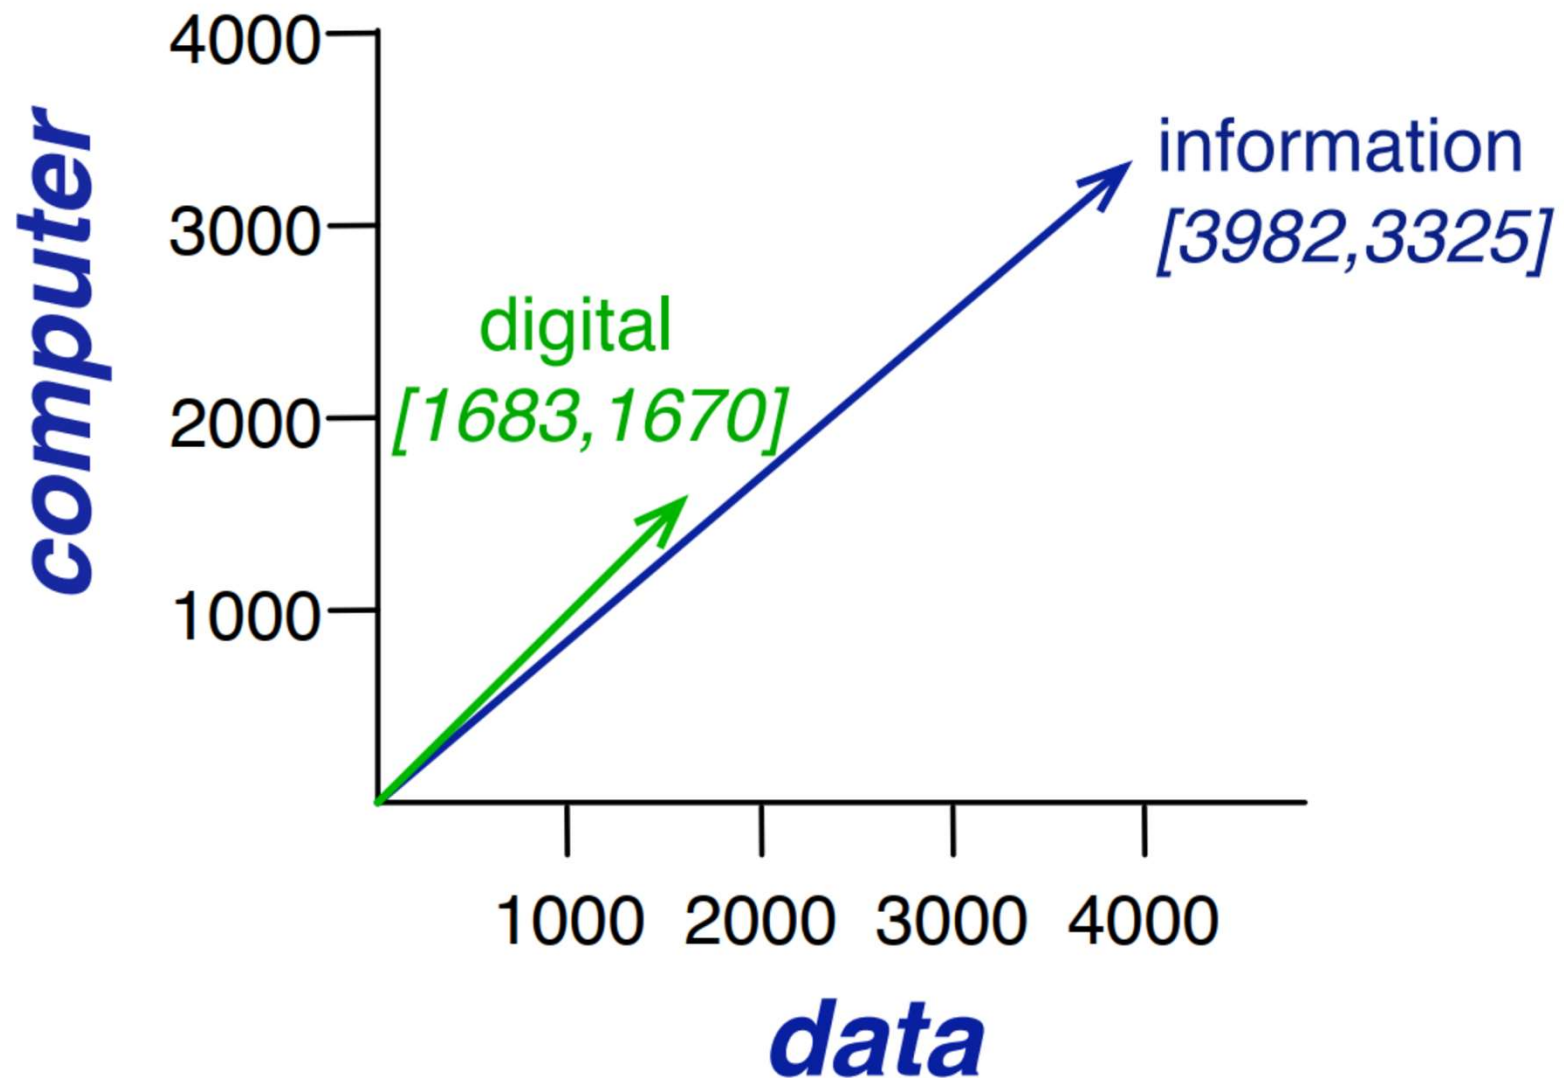
Word Similarity

- Two words are **similar in meaning** if their **context vectors are similar**

is traditionally followed by **cherry** pie, a traditional dessert
often mixed, such as **strawberry** rhubarb pie. Apple pie
computer peripherals and personal **digital** assistants. These devices usually
a computer. This includes **information** available on the internet

	aardvark	...	computer	data	result	pie	sugar	...
cherry	0	...	2	8	9	442	25	...
strawberry	0	...	0	0	1	60	19	...
digital	0	...	1670	1683	85	5	4	...
information	0	...	3325	3982	378	5	13	...

Word Similarity Visualization



Word Similarity | Cosine Similarity

$$\cos(\vec{v}, \vec{w}) = \frac{\vec{v} \cdot \vec{w}}{|\vec{v}| |\vec{w}|} = \frac{\vec{v} \cdot \vec{w}}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}} = \frac{\sum_{i=1}^N v_i w_i}{\sqrt{\sum_{i=1}^N v_i^2} \sqrt{\sum_{i=1}^N w_i^2}}$$

	pie	data	computer
cherry	442	8	2
digital	5	1683	1670
information	5	3982	3325

$$\cos(\text{cherry}, \text{information}) =$$

$$\frac{442 * 5 + 8 * 3982 + 2 * 3325}{\sqrt{442^2 + 8^2 + 2^2} \sqrt{5^2 + 3982^2 + 3325^2}} = .017$$

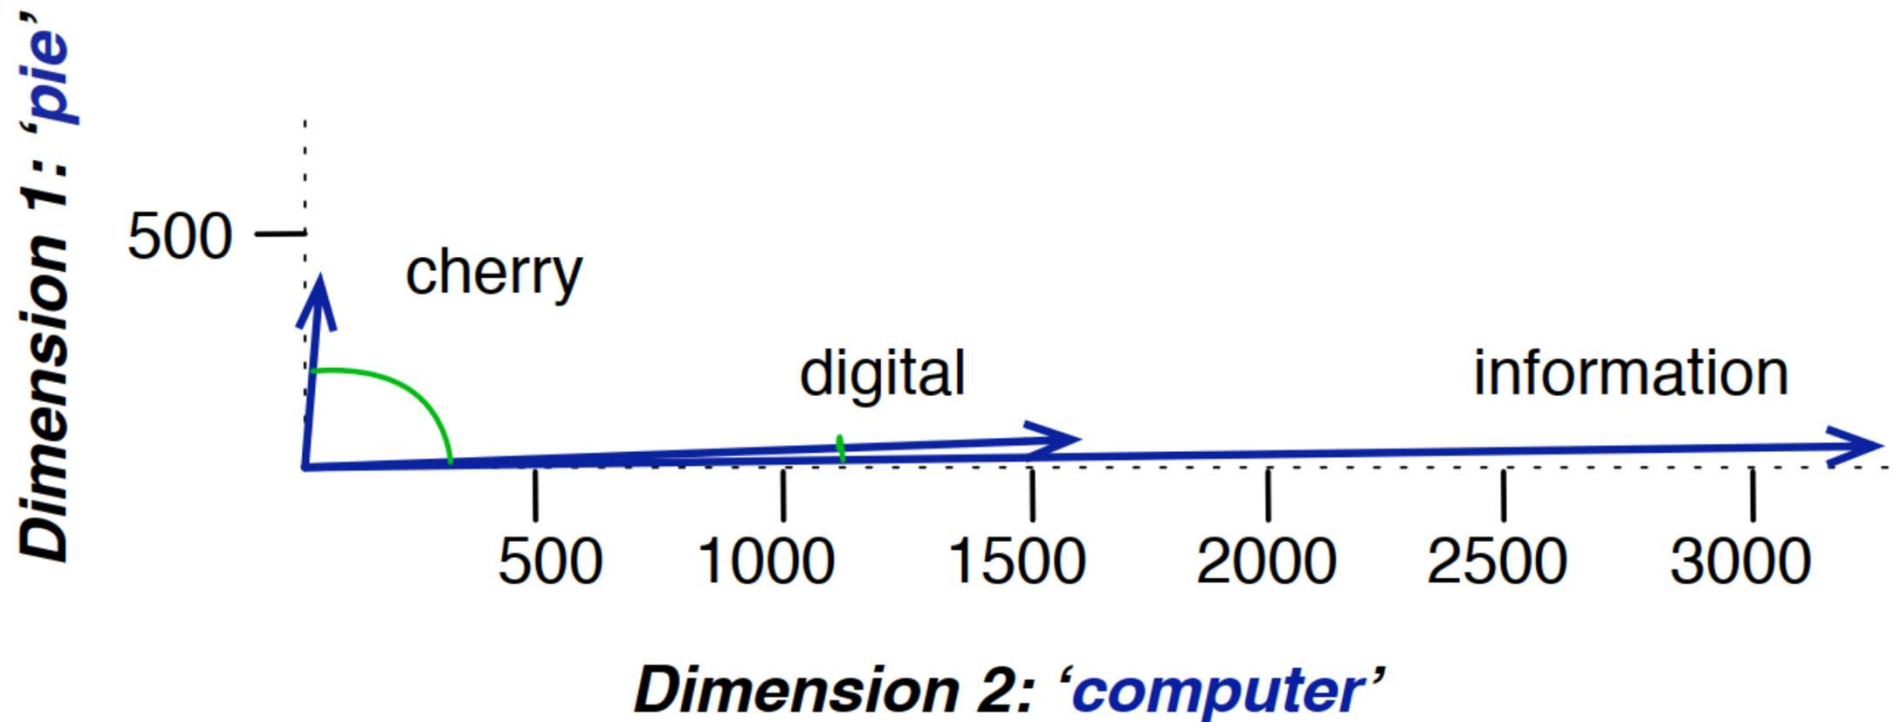
$$\cos(\text{digital}, \text{information}) =$$

Low similarity

$$\frac{5 * 5 + 1683 * 3982 + 1670 * 3325}{\sqrt{5^2 + 1683^2 + 1670^2} \sqrt{5^2 + 3982^2 + 3325^2}} = .996$$

High similarity

Cosine Similarity Visualization



Vector Embeddings: Methods

- `tf-idf`
 - popular in Information Retrieval
 - **sparse** vectors
 - word represented by **(a simple function of) the counts of nearby words**
- `Word2vec`
 - **dense** vectors
 - representation is created by training **a classifier to predict whether a word is likely to appear nearby**

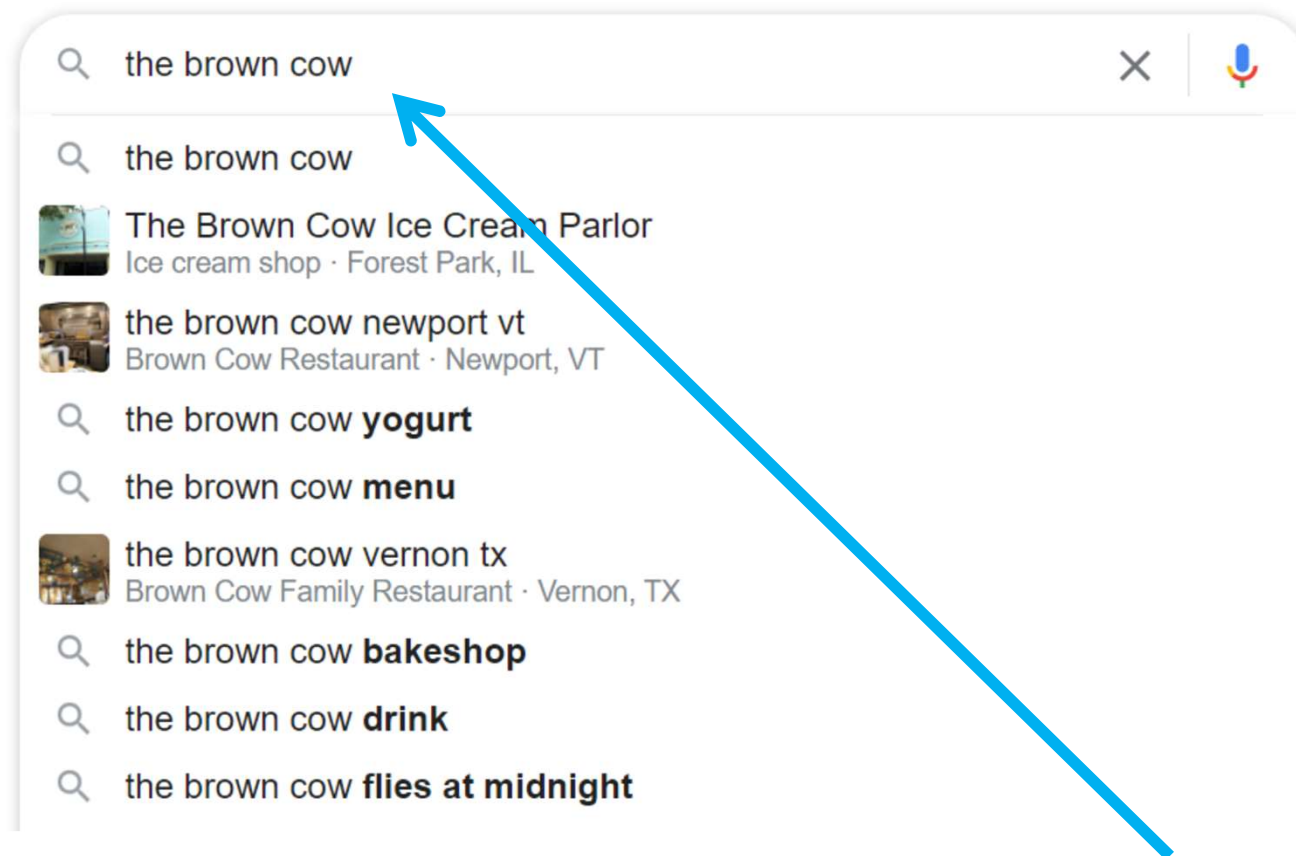
Sparse vs. Dense Vectors

- **Sparse vectors have a lot of values set to zero.**
- **Dense vector: most of the values are non-zero.**
 - **better use of storage**
 - **carries more information**

tf-idf: Frequencies Not Enough

- The co-occurrence matrices we have seen represent each cell by **word frequencies**
- Frequency is clearly useful:
 - if *sugar* appears a lot **near** *apricot*, that's **useful information**
- But overly frequent words like *the*, *it*, or *they* are **not very informative** about the context
- It's a paradox! How can we balance these two conflicting constraints?

tf-idf: Motivation



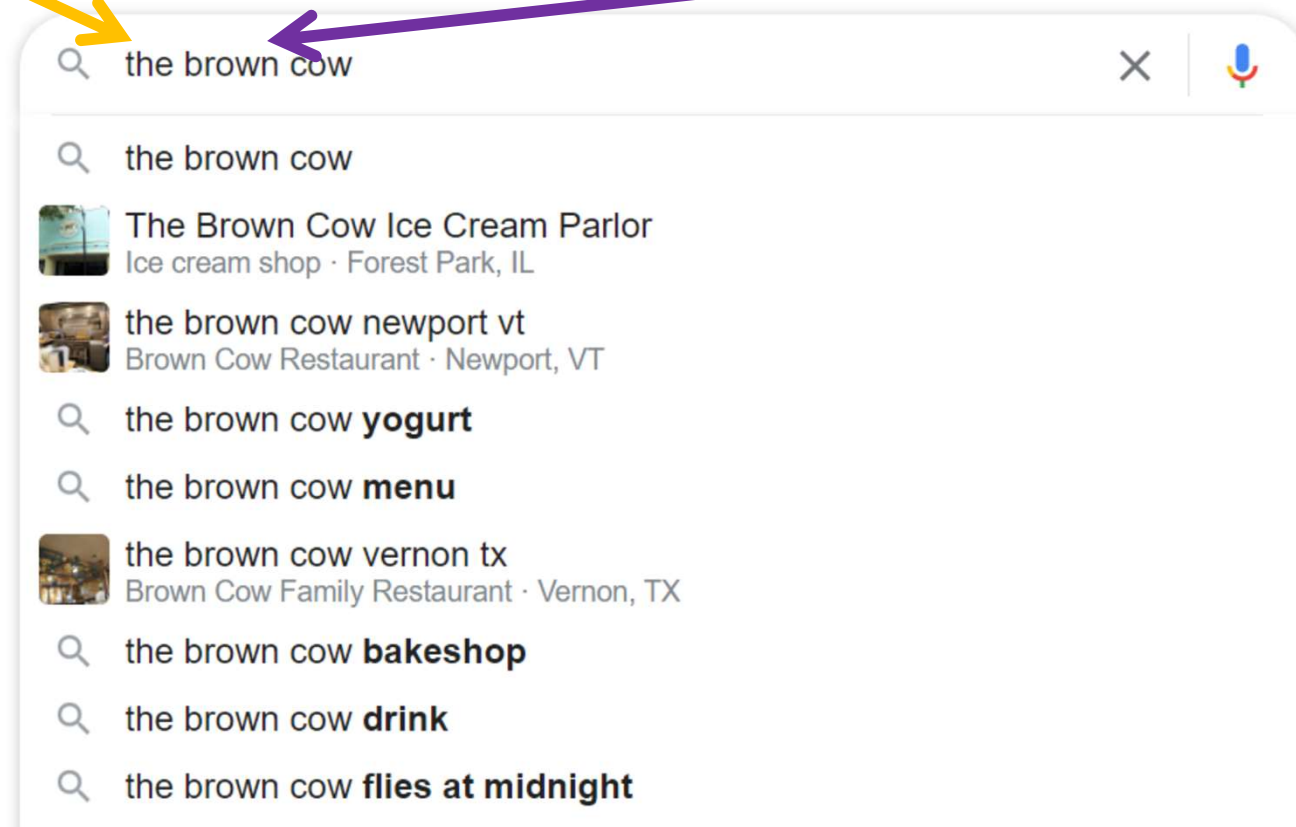
we want to find **documents** most **relevant** to the **search term**

tf-idf: Motivation

“the” is a **very common word** =
high frequency



“brown” and “cow” are
less common words =
lower frequency



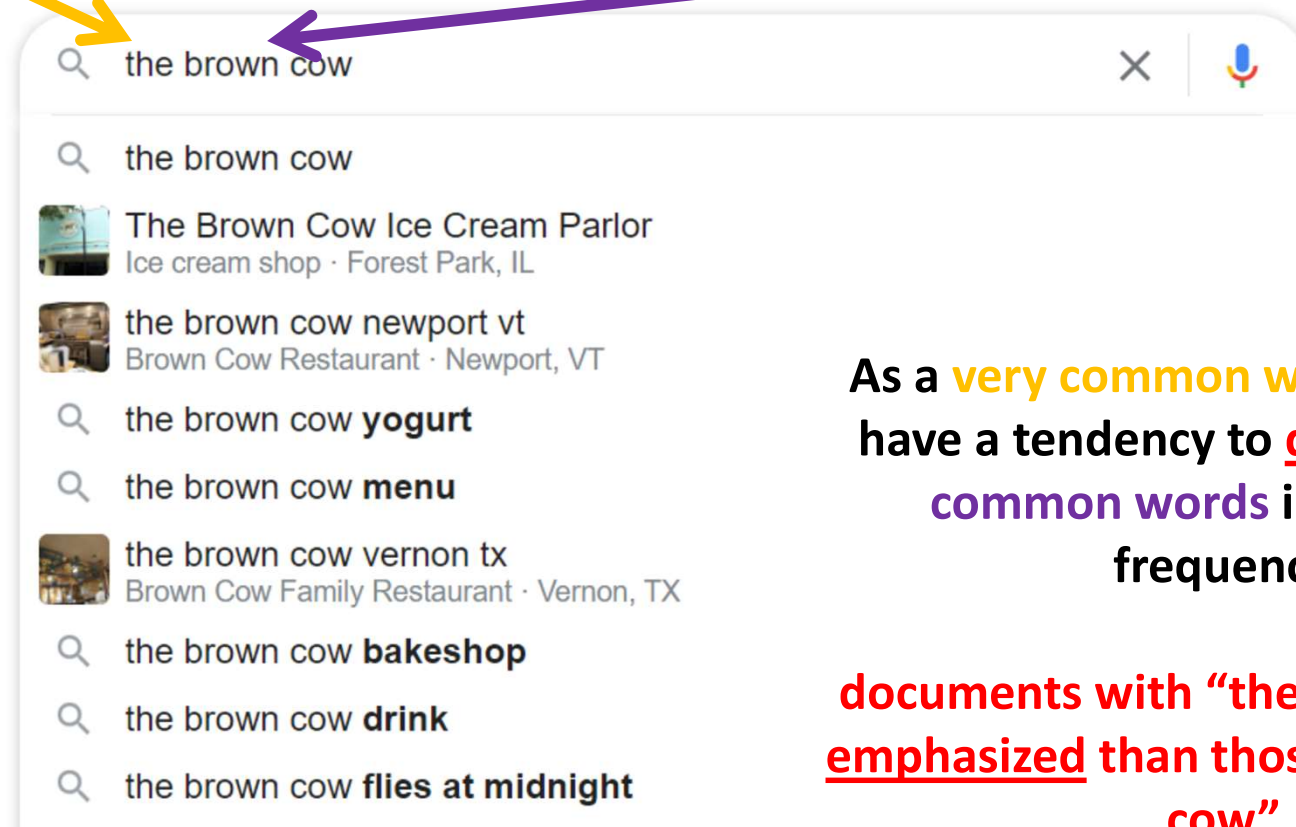
we want to find **documents** most **relevant** to the **search term**

tf-idf: Motivation

“the” is a **very common word** =
high frequency



“brown” and “cow” are
less common words =
lower frequency



As a **very common word** “the” will
have a tendency to **dominate** **less
common words** in terms of
frequency

**documents with “the” will be more
emphasized than those with “brown
cow”**

we want to find **documents** most **relevant** to the **search term**

tf-idf: Motivation

- **Problem 1: Vectorization can be “not representative”**
 - binary vectorization gives too little weight to words that occur multiple times.
 - count vectorization gives too much weight.
 - A single word that shows up many times can **drown** the effect of the rest of the text in determining what other texts are most similar
- **Problem 2: Not all words should be equal**
 - words with **distinctive meaning** (*turbine, diagonalize, cantilever*) **should count more** towards the document representation than very general terms (*people, things, stuff, day*)

tf-idf

term **f**requency – **i**nverse **d**ocument **f**requency

tf-idf: Idea and Intuition

- Idea:

- define a “score” for each component of the document vector associated with a given word w , broken down into two parts:
 - some measure of the frequency of w within the individual document (tf)
 - some measure of the rarity of w across all documents (corpus)(idf)

- Intuition:

- tf can be chosen so that a lot of occurrences doesn't add up to too high a score (to address Problem 1),
- idf will penalize the score for words that show up all over the place (to address Problem 2)

tf-idf

term frequency – inverse document frequency



how often does the term
(word) appear in a specific
document?



1 / how often does the term (word) appear in the
entire corpus?

tf-idf

term frequency – inverse document frequency



how often does the term (word) appear in a specific document?



how distinctive is the term (word) for a specific document?



1 / how often does the term (word) appear in the entire corpus?



1 / how common / popular is the term (word) within the entire corpus?

tf-idf

term frequency – inverse document frequency



how often does the term (word) appear in a specific document?



how distinctive is the term (word) for a specific document?



1 / how often does the term (word) appear in the entire corpus?



1 / how common / popular is the term (word) within the entire corpus?

we want to balance out term frequency with inverse document frequency

tf-idf

term frequency – inverse document frequency

$$\text{tf}(\text{word}, \text{document } d_i) =$$

$$= \text{count} \left(\begin{array}{c} \text{document } d_i \\ \begin{array}{|c|} \hline \begin{array}{c} \dots \text{word} \dots \\ \text{word} \dots \\ \dots \text{word} \\ \dots \text{word} \dots \\ \text{word} \dots \end{array} \\ \hline \end{array} \right) \end{array} \right)$$

$$\text{idf}(\text{word}, \text{corpus } C) =$$

$$= \log \left(\frac{N: \text{number of all documents in } C}{\text{number of all documents with word}} \right)$$

tf-idf

term frequency – inverse document frequency

$$\text{tf}(\text{word}, \text{document } d_i) =$$

$$= \text{count} \left(\begin{array}{c} \text{document } d_i \\ \begin{array}{|c|} \hline \begin{array}{c} \dots \text{word} \dots \\ \text{word} \dots \\ \dots \text{word} \\ \dots \text{word} \dots \\ \text{word} \dots \end{array} \\ \hline \end{array} \right) \end{array} \right)$$

$$\text{idf}(\text{word}, \text{corpus } C) =$$

$$= \log \left(\frac{N: \text{number of all documents in } C}{\text{number of all documents including word}} \right)$$

$$\text{tfidf}(\text{word}, \text{document } d_i, \text{corpus } C) = \text{tf}(\text{word}, \text{document } d_i) * \text{idf}(\text{word}, \text{corpus } C) =$$

$$= \text{count}(\text{word in document } d_i) * \log \left(\frac{N: \text{number of all documents in } C}{\text{number of all documents including word}} \right)$$

tf-idf

term frequency – inverse document frequency

$$\begin{aligned}\text{tfidf}(\text{word}, \text{document } d_i, \text{corpus } C) &= \text{tf}(\text{word}, \text{document } d_i) * \text{idf}(\text{word}, \text{corpus } C) = \\ &= \text{count}(\text{word} \text{ in document } d_i) * \log\left(\frac{N: \text{number of all documents in } C}{\text{number of all documents including word}}\right)\end{aligned}$$

$\text{tfidf}(\text{word}, \text{document } d_i, \text{corpus } C)$ goes \uparrow
for words that are **very specific** to document d_i but **not common**
in corpus C

$\text{tfidf}(\text{word}, \text{document } d_i, \text{corpus } C)$ goes \downarrow
for words that are **NOT very specific** to document d_i but **common**
in corpus C

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

$$\text{tfidf}(\text{"this"}, \text{document } d_1, \text{corpus } C) = \text{tf}(\text{"this"}, \text{document } d_1) * \text{idf}(\text{"this"}, \text{corpus } C) =$$

$$= \text{count}(\text{"this"} \text{ in document } d_1) * \log\left(\frac{N: \text{number of all documents in } C}{\text{number of all documents including "this"}}\right)$$

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

$$\text{tfidf}(\text{"this"}, \text{document } d_1, \text{corpus } C) = \text{tf}(\text{"this"}, \text{document } d_1) * \text{idf}(\text{"this"}, \text{corpus } C) =$$

$$= 1 * \log\left(\frac{2}{2}\right) = 1 * 0 = 0$$

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

$\text{tfidf}(\text{"sample"}, \text{document } d_1, \text{corpus } C) = \text{tf}(\text{"sample"}, \text{document } d_1) * \text{idf}(\text{"sample"}, \text{corpus } C) =$

$= \text{count}(\text{"sample"} \text{ in document } d_1) * \log\left(\frac{N: \text{number of all documents in } C}{\text{number of all documents including "sample"}}\right)$

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

$$\begin{aligned}\text{tfidf}(\text{"sample"}, \text{document } d_1, \text{corpus C}) &= \text{tf}(\text{"sample"}, \text{document } d_1) * \text{idf}(\text{"sample"}, \text{corpus C}) = \\ &= 1 * \log\left(\frac{2}{1}\right) = 1 * 1 = 1\end{aligned}$$

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

$$\text{tfidf}(\text{"is"}, \text{document } d_1, \text{corpus } C) = \text{tf}(\text{"is"}, \text{document } d_1) * \text{idf}(\text{"is"}, \text{corpus } C) =$$

$$= \text{count}(\text{"is" in document } d_1) * \log\left(\frac{N: \text{number of all documents in } C}{\text{number of all documents including "is"}}\right)$$

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

$$\text{tfidf}(\text{"is"}, \text{document } d_1, \text{corpus C}) = \text{tf}(\text{"is"}, \text{document } d_1) * \text{idf}(\text{"is"}, \text{corpus C}) =$$

$$= 1 * \log\left(\frac{2}{2}\right) = 1 * 0 = 0$$

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

$$\begin{aligned}
 \text{tfidf}(\text{"a"}, \text{document } d_1, \text{corpus } C) &= \text{tf}(\text{"a"}, \text{document } d_1) * \text{idf}(\text{"a"}, \text{corpus } C) = \\
 &= \text{count}(\text{"a"} \text{ in document } d_1) * \log\left(\frac{N: \text{number of all documents in } C}{\text{number of all documents including "a"}}\right)
 \end{aligned}$$

tf-idf: Example

document d_1

this is a
sample

Corpus C

document d_1

this is a
sample

document d_2

this is
another
example

word	count
this	1
is	1
a	1
sample	1

word	count	word	count
this	2	another	1
is	2	example	1
a	1		
sample	1		

$$\begin{aligned}\text{tfidf}(\text{"a"}, \text{document } d_1, \text{corpus C}) &= \text{tf}(\text{"a"}, \text{document } d_1) * \text{idf}(\text{"a"}, \text{corpus C}) = \\ &= 1 * \log\left(\frac{2}{1}\right) = 1 * 1 = 1\end{aligned}$$

tf-idf: Example

document d_1

this is a
sample

word	count
this	1
is	1
a	1
sample	1

tf-idf



word	tfidf
this	0
is	0
a	1
sample	1

tf-idf: Alternative Measures

term **f**requency - **i**nverse **d**ocument **f**requency

scheme	tf weight
binary	0, 1
raw count (used in example)	$\text{count}(\text{word in document})$
term frequency	$\frac{\text{count}(\text{word in document})}{\sum_i \text{count}(\text{anyWord}_i \text{ in document})}$
log normalization	$\log(1 - \text{count}(\text{word in document}))$

scheme	idf weight
unary	1
inverse document frequency (used in example)	$\log\left(\frac{N: \text{number of all documents in Corpus}}{\text{number of all documents including word}}\right)$ $= -\log\left(\frac{\text{number of all documents including word}}{N: \text{number of all documents in Corpus}}\right)$
inverse document frequency smooth	$1 + \log\left(\frac{N: \text{number of all documents in Corpus}}{1 + \text{number of all documents including word}}\right)$

tf-idf: Alternative Measures

term frequency - **inverse document frequency**

scheme	tf weight
binary	0, 1
raw count	$\text{count}(\text{word in document})$
term frequency	$\frac{\text{count}(\text{word in document})}{\sum_i \text{count}(\text{anyWord}_i \text{ in document})}$
log normalization (typical)	$\log(1 - \text{count}(\text{word in document}))$

scheme	idf weight
unary	1
inverse document frequency (typical)	$\log\left(\frac{N: \text{number of all documents in Corpus}}{\text{number of all documents including word}}\right)$ $= -\log\left(\frac{\text{number of all documents including word}}{N: \text{number of all documents in Corpus}}\right)$
inverse document frequency smooth	$1 + \log\left(\frac{N: \text{number of all documents in Corpus}}{1 + \text{number of all documents including word}}\right)$

Words as Vectors: Issues

- We saw how to build vectors to represent words:
 - one-hot encoding:
 - binary, count, $tf*idf$
- Some problems
 - Large dimensionality of word vectors
 - Lack of meaningful relationships between words

Vector Embeddings: Methods

- `tf-idf`
 - popular in Information Retrieval
 - **sparse** vectors
 - word represented by **(a simple function of) the counts of nearby words**
- `Word2vec`
 - **dense** vectors
 - representation is created by training **a classifier to predict whether a word is likely to appear nearby**

Sparse vs. Dense Vectors

- **Sparse vectors have a lot of values set to zero.**

[0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 2]

- **Dense vector: most of the values are non-zero.**

- **better use of storage**
- **carries more information**

[3, 1, 5, 0, 1, 4, 9, 8, 7, 1, 1, 2, 2, 2, 2]

Sparse vs. Dense Vectors

- `tf-idf` vectors are typically:
 - **long** (length 20,000 to 50,000)
 - **sparse** (most elements are zero)
- What if we could learn vectors that are
 - **short** (length 50-1000)
 - **dense** (most elements are non-zero)

Short / Dense Vectors: Benefits

- Why short/dense vectors?
 - short vectors may be easier to use as **features** in machine learning (fewer weights to tune)
 - dense vectors may **generalize** better than explicit counts
 - dense vectors may do better at capturing synonymy:
 - *car* and *automobile* are synonyms; but are distinct dimensions
 - a word with *car* as a neighbor and a word with *automobile* as a neighbor should be similar, but aren't
- **In practice, they work better**

Short/Dense Vectors: Methods

- “Neural Language Model”-inspired models
 - **Word2vec**, GloVe
- Singular Value Decomposition (SVD)
 - A special case of this is called LSA – Latent Semantic Analysis
- Alternative to these “static embeddings”:
 - Contextual Embeddings (ELMo, BERT)
 - Compute distinct embeddings for a word in its context
 - Separate embeddings for each token of a word

Word2Vec: Idea

DON'T count - Predict!

Word2Vec: Idea

- Instead of **counting** how often each **word** w occurs near "*apricot*"
 - Train a classifier on a binary **prediction** task:
 - Is w likely to show up near "*apricot*"?
- We don't actually care about this task
 - but we'll take the learned classifier weights as the word embeddings
- Use **self-supervision**:
 - A **word** c that occurs near "*apricot*" in the corpus acts as the gold "correct answer" for supervised learning
 - **No need for human labels**

Available Tools

- Word2vec (Mikolov et al)

<https://code.google.com/archive/p/word2vec/>

- GloVe (Pennington, Socher, Manning)

<http://nlp.stanford.edu/projects/glove/>

Word2Vec: the Approach

1. Treat the target **word** t and a neighboring context **word** c as **positive examples**.
2. Randomly sample other words in the lexicon to get **negative examples**
3. Use **logistic regression** to train a classifier to distinguish those two cases
4. Use the **learned classifier weights** as the **embeddings**

Word2Vec: the Approach

Given the set of **positive** and **negative** training instances, and an **initial set of embedding vectors**

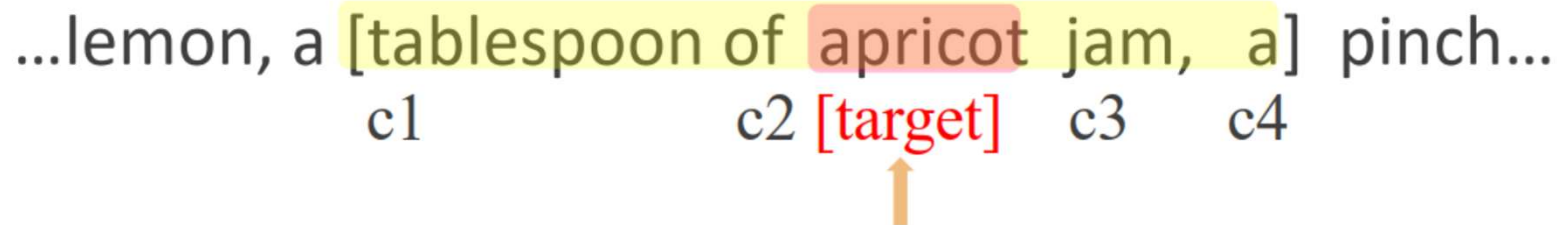
The goal of learning is to adjust those word vectors such that we:

- **maximize** the similarity of the **target word**, **context word** pairs (w , c_{pos}) drawn from the **positive** data
- **minimize** the similarity of the (w , c_{neg}) pairs drawn from the **negative** data.

Word2Vec: the Approach

...lemon, a [tablespoon of apricot jam, a] pinch...

c1 c2 [target] c3 c4



positive examples +

t	c
apricot	tablespoon
apricot	of
apricot	jam
apricot	a

For each positive example we'll grab k negative examples, sampling by frequency

Word2Vec: the Approach

...lemon, a [tablespoon of apricot jam, a] pinch...

c1 c2 [target] c3 c4

↑

positive examples +

t	c
apricot	tablespoon
apricot	of
apricot	jam
apricot	a

negative examples -

t	c	t	c
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if